



Q1: TensorFlow vs. PyTorch

TensorFlow:

Static computation graph (define-then-run), optimized for deployment.

Strong production tools (TF Serving, Lite).

Ideal for large-scale systems and mobile deployment.

PyTorch:

Dynamic computation graph (define-by-run), easier debugging.

Pythonic syntax, preferred for research/prototyping.

When to choose:

TensorFlow: Production pipelines, edge devices.

PyTorch: Rapid experimentation, academia.

Q2: Jupyter Notebook Use Cases

Interactive Prototyping:

Visualize data/model outputs (e.g., matplotlib plots) during exploratory analysis.

Educational Demos:

Combine code, equations, and explanations (Markdown) to teach ML concepts.

Q3: spaCy vs. Basic String Operations

spaCy:

Pre-trained models for POS tagging, dependency parsing, and NER.

Handles linguistic nuances (e.g., "Apple" as company vs. fruit).

Basic Strings:

Limited to regex/string matching; fails at context (e.g., irony in "This 'gift' ruined my day").

Comparative Analysis: Scikit-learn vs. TensorFlow

Scikit-learn:

- | | |
|-----------------------------|--|
| Target Applications: | Classical ML (SVM, decision trees) |
| Ease of Use: | Simple API, less code for traditional ML |
| Community Support: | Extensive docs/tutorials for ML basics |

TensorFlow:

- | | |
|-----------------------------|---|
| Target Applications: | Deep Learning (CNNs, RNNs) |
| Ease of Use: | Steeper learning curve; flexible for DL |
| Community Support: | Massive ecosystem; industry resources |